

gaussian_renderer directory and GaussianRasterizer class

- ./gaussian_renderer/
- render() method

```
# gaussian_renderer/__init__.py
...
from diff_gaussian_rasterization import GaussianRasterizationSettings,
GaussianRasterizer
...
def render(viewpoint_camera, pc : GaussianModel, pipe, bg_color : torch.Tensor,
scaling_modifier = 1.0, override_color = None):
    ...
    tanfovx = math.tan(viewpoint_camera.FoVx * 0.5)
    tanfovy = math.tan(viewpoint_camera.FoVy * 0.5)

    raster_settings = GaussianRasterizationSettings(
        image_height=int(viewpoint_camera.image_height),
        image_width=int(viewpoint_camera.image_width),
        tanfovx=tanfovx,
        tanfovy=tanfovy,
        bg=bg_color,
        scale_modifier=scaling_modifier,
        viewmatrix=viewpoint_camera.world_view_transform,
        projmatrix=viewpoint_camera.full_proj_transform,
        sh_degree=pc.active_sh_degree,
        campos=viewpoint_camera.camera_center,
        prefiltered=False,
        debug=pipe.debug
    )

    rasterizer = GaussianRasterizer(raster_settings=raster_settings)

    ...
    rendered_image, radii = rasterizer(
        means3D = means3D,
        means2D = means2D,
        shs = shs,
        colors_precomp = colors_precomp,
        opacities = opacity,
        scales = scales,
        rotations = rotations,
        cov3D_precomp = cov3D_precomp)

    return {"render": rendered_image,
            "viewspace_points": screenspace_points,
            "visibility_filter" : radii > 0,
            "radii": radii}
```

- **Rendering:** Rendering is the broader process of generating an image from a model. This model can include geometry (shapes and structures), texture (surface detail), lighting (how light interacts with surfaces), and shading (how colors and shadows are applied).
- **Rasterization:** **Rasterization is a specific method used within the rendering process** to convert 3D models into 2D images. It translates geometric data (like vertices and edges) into pixels or fragments.

Parameters for Rendering

- **viewpoint_camera:** the rendered image will be based on the viewpoint_camera
- **pc (point cloud):** 3D gaussian points that will be **splatted** on to the viewpoint camera plane
- **pipe:** pipe is created as the training starts. The pipe contains attributes to determine whether certain operation should be performed at certain place, e.g.,:

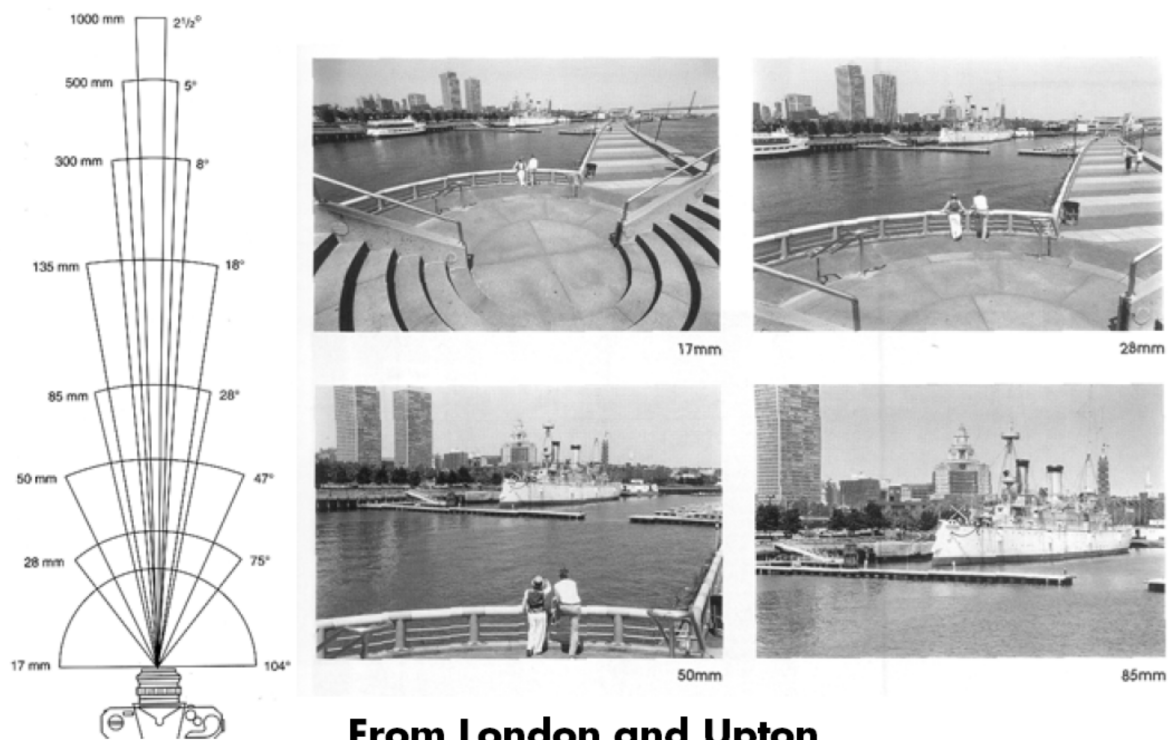
```
# whether need to precompute 3D covariance
if pipe.compute_cov3D_python:
    cov3D_precomp = pc.get_covariance(scaling_modifier)
# whether need to convert SH to rgb before rasterization
if pipe.convert_SHs_python:
    shs_view = pc.get_features.transpose(1, 2).view(-1, 3,
    (pc.max_sh_degree+1)**2)
    dir_pp = (pc.get_xyz -
    viewpoint_camera.camera_center.repeat(pc.get_features.shape[0], 1))
    dir_pp_normalized = dir_pp/dir_pp.norm(dim=1, keepdim=True)
    sh2rgb = eval_sh(pc.active_sh_degree, shs_view, dir_pp_normalized)
    colors_precomp = torch.clamp_min(sh2rgb + 0.5, 0.0)
```

- **bg_color:** background color. Will occupy pixels that not have been rendered
- **scaling_modifier:** in the case that the resolution of the image is too large, may use scaling_modifier > 1 to downsample, similar to resolution_scales
- **override_color:** maybe a feature that can change the original color that whatever color we want

Parameters for GaussianRasterizationSettings

- **image_height:** height of the rendered image
- **image_width:** width of the rendered image
- **tanfovx:** tangent of field of view in x-axis
 - **fov :** field of view(视场角\视野), quantified by angle

Field of View (Zoom)



From London and Upton

- tanfovy: tangent of field of view in y-axis
- bg: background color
- scale_modifier: in the case that the resolution of the image is too large, may use scaling_modifier > 1 to downsample, similar to resolution_scales
- viewmatrix: viewpoint_camera.world_view_transform, the matrix defines the relation between the world space (3D space) and the camera space (view space)
- projmatrix: viewpoint_camera.full_proj_transform, combining world_view_matrix and projection matrix
 - relevant code:

```
# scene/cameras.py
self.world_view_transform = torch.tensor(getWorld2View2(R, T, trans,
scale)).transpose(0, 1).cuda()
self.projection_matrix = getProjectionMatrix(znear=self.znear,
zfar=self.zfar, fovX=self.FoVx, fovY=self.FoVy).transpose(0,1).cuda()
self.full_proj_transform =
(self.world_view_transform.unsqueeze(0).bmm(self.projection_matrix.unsqueeze
(0))).squeeze(0)
self.camera_center = self.world_view_transform.inverse()[3, :3]

# utils/graphics_utils.py
def getWorld2View2(R, t, translate=np.array([.0, .0, .0]), scale=1.0):
    Rt = np.zeros((4, 4))
    Rt[:3, :3] = R.transpose()
    Rt[:3, 3] = t
    Rt[3, 3] = 1.0
```

```

C2W = np.linalg.inv(Rt)
cam_center = C2W[:3, 3]
cam_center = (cam_center + translate) * scale
C2W[:3, 3] = cam_center
Rt = np.linalg.inv(C2W)
return np.float32(Rt)

def getProjectionMatrix(znear, zfar, fovX, fovY):
    tanHalfFovY = math.tan((fovY / 2))
    tanHalfFovX = math.tan((fovX / 2))

    top = tanHalfFovY * znear
    bottom = -top
    right = tanHalfFovX * znear
    left = -right

    P = torch.zeros(4, 4)

    z_sign = 1.0

    P[0, 0] = 2.0 * znear / (right - left)
    P[1, 1] = 2.0 * znear / (top - bottom)
    P[0, 2] = (right + left) / (right - left)
    P[1, 2] = (top + bottom) / (top - bottom)
    P[3, 2] = z_sign
    P[2, 2] = z_sign * zfar / (zfar - znear)
    P[2, 3] = -(zfar * znear) / (zfar - znear)
    return P

```

- campos: the position of the camera (center of the viewpoint)
- prefiltered: refer to the low-pass filter in EWA splatting
 - relevant code: (this is a cuda function determine if a point in the screen. It seems that `prefiltered` indicates that the points have undergone some preprocessing, filtering points with `prefiltered=True` will trigger some warnings)
 - EWA splatting: We can either sample the continuous signal at a higher frequency or we eliminate frequencies above the Nyquist limit before sampling, which is called prefiltering.

```

__forceinline__ __device__ bool in_frustum(int idx,
const float* orig_points,
const float* viewmatrix,
const float* projmatrix,
bool prefiltered,
float3& p_view)
{
    float3 p_orig = { orig_points[3 * idx], orig_points[3 * idx + 1],
orig_points[3 * idx + 2] };

    // Bring points to screen space
    float4 p_hom = transformPoint4x4(p_orig, projmatrix);
    float p_w = 1.0f / (p_hom.w + 0.0000001f);
    float3 p_proj = { p_hom.x * p_w, p_hom.y * p_w, p_hom.z * p_w };

```

```

p_view = transformPoint4x3(p_orig, viewmatrix);

if (p_view.z <= 0.2f) // || ((p_proj.x < -1.3 || p_proj.x > 1.3 ||
p_proj.y < -1.3 || p_proj.y > 1.3)))
{
    if (prefiltered)
    {
        printf("Point is filtered although prefiltered is set. This
shouldn't happen!");
        __trap();
    }
    return false;
}
return true;
}

```

- debug: whether enable debug setting

Many of the parameters are related to cameras

- an example camera setting:

```

[{"id": 0, "img_name": "DSC05572", "width": 1264, "height": 832, "position":
[3.1404339644832397, 0.18188197960281, -3.563482533678278], "rotation":
[[0.9979604943324005, 0.011462677284412257, -0.06279696474594969],
[-0.04072327873048568, 0.8718824481262268, -0.4880190684992303],
[0.049157566266748824, 0.4895810491419074, 0.8705710367338437]], "fy":
1040.192756666184, "fx": 1040.007303759328}]

```

- cameras will be loaded by `utils/camera_utils.py` and then initialized by `scene/cameras.py`

submodules\diff-gaussian-rasterization

- Differential Gaussian Rasterization
- A package contains functions that perform rasterization (rendering)
- Most of the parameters above will be used in this package to perform rasterization

Relation to EWA splatting

- EWA: the whole object model is a resampling of a continuous function
- 3DGS: the point cloud is a resampling of a continuous function
- Volume Resampling in EWA & Rasterization in 3DGS
 - two ways of volume rendering:
 - backward mapping (NeRF)
 - forward mapping (EWA & 3DGS)
 - EWA: Mapping the data onto the image plane involves a sequence of intermediate steps where the data is transformed to different coordinate systems
 - final goal: 3D points -> 2D images that can be perceived by human

Data Flow

- source space / object space / world space -> camera space -> ray space -> screen space (viewport coordinates)

Detailed Implementation

- `submodules\diff-gaussian-rasterization\diff_gaussian_rasterization__init__.py` invokes C++/CUDA rasterizer
 - examples of invoking a C++/CUDA function:

```

from . import _C
...
# Invoke C++/CUDA rasterizer
if raster_settings.debug:
    cpu_args = cpu_deep_copy_tuple(args) # Copy them before they can be
    corrupted
    try:
        num_rendered, color, radii, geomBuffer, binningBuffer, imgBuffer
    = _C.rasterize_gaussians(*args)
    except Exception as ex:
        torch.save(cpu_args, "snapshot_fw.dump")
        print("\nAn error occured in forward. Please forward
snapshot_fw.dump for debugging.")
        raise ex
    else:
        num_rendered, color, radii, geomBuffer, binningBuffer, imgBuffer =
_C.rasterize_gaussians(*args)
...
# Compute gradients for relevant tensors by invoking backward method
if raster_settings.debug:
    cpu_args = cpu_deep_copy_tuple(args) # Copy them before they can be
    corrupted
    try:
        grad_means2D, grad_colors_precomp, grad_opacities, grad_means3D,
grad_cov3Ds_precomp, grad_sh, grad_scales, grad_rotations =
_C.rasterize_gaussians_backward(*args)
    except Exception as ex:
        torch.save(cpu_args, "snapshot_bw.dump")
        print("\nAn error occured in backward. Writing snapshot_bw.dump for
debugging.\n")
        raise ex
    else:
        grad_means2D, grad_colors_precomp, grad_opacities, grad_means3D,
grad_cov3Ds_precomp, grad_sh, grad_scales, grad_rotations =
_C.rasterize_gaussians_backward(*args)
...
def markVisible(self, positions):
    # Mark visible points (based on frustum culling for camera) with a
    boolean
    with torch.no_grad():
        raster_settings = self.raster_settings

```

```

        visible = _C.mark_visible(
            positions,
            raster_settings.viewmatrix,
            raster_settings.projmatrix)

    return visible

```

C++/CUDA rasterizer details

- directory hierachy:

```

diff-gaussian-rasterization/
├── cuda_rasterizer/          # C++/CUDA rasterizer implementation
│   ├── auxiliary.h          # Auxiliary functions and definitions
│   ├── backward.cu          # CUDA kernel for backward pass (gradient computation)
│   ├── backward.h           # Header file for backward pass
│   ├── forward.cu           # CUDA kernel for forward pass (rasterization)
│   ├── forward.h            # Header file for forward pass
│   ├── rasterizer_impl.cu   # Main CUDA rasterizer implementation
│   ├── rasterizer_impl.h    # Header file for main rasterizer implementation
│   └── rasterizer.h         # Header file for rasterizer interface
├── diff_gaussian_rasterization/ # Python initialization module
│   ├── __init__.py          # Python package initialization
├── third_party/              # Third-party packages
│   └── glm/                  # OpenGL Mathematics library
├── CMakeLists.txt            # CMake configuration file for building the project
├── ext.cpp                   # C++ extension module definition for Python
├── rasterize_points.cu       # CUDA kernel for rasterizing points
├── rasterize_points.h        # Header file for point rasterization
└── setup.py                  # Python setup script for building and installing the
package

```

- forward process (forward.cu & forward.h):
 - forward mapping mentioned in EWA splatting
 - the process of rendering and rasterization
 - obtain rendered image that can compute loss
- backward process (backward.cu & backward.h):
 - computing gradient for minimizing loss
 - optimize the parameters

EWA Splatting

Abstract

- high quality splatting using elliptical Gaussian kernel
- avoid aliasing artifacts (can refer to [\[常见的伪影\]](#))
 - image aliasing:

- Moire Pattern (摩尔纹)
- Jagged Edge (锯齿状边缘)
- Pixelation (像素化)
- video aliasing:
 - Temporal Aliasing: frame rate is too low
 - Flickering (闪烁效应)
- introduce the concept of a resampling filter, combining a reconstruction kernel with a low-pass filter
- like EWA (elliptical weighted average) filter for texture mapping -> EWA splatting
- EWA splat primitives
- can be used in regular, rectilinear, and irregular dataset
- EWA volume reconstruction kernels can be reduced to surface reconstruction kernels

Introduction

Ideal Volume Rendering

- The ideal volume rendering algorithm **reconstructs a continuous function** in 3D, **transforms** this 3D function into screen space, and then **evaluates opacity integrals** along line-of-sights
 - Reconstructing a continuous function in 3D
 - The 3D data is often stored as a discrete set of samples (voxels). Reconstructing a continuous function means interpolating these discrete samples to create a smooth, continuous representation of the 3D volume
 - may use interpolation
 - Transforms to screen space
 - project the function to 2D space so that human can see
 - Evaluates opacity integrals
 - Imagine casting rays from the viewer's eye through each pixel on the screen into the 3D volume. Each ray represents a line of sight

Splatting Algorithm

- volume data -> set of particles absorbing and emitting light
- line integrals are precomputed across each particles separately -> **footprint functions / splats** -> final image
- elliptical weighted average texture filter
- concept of ray space:
 - a point in ray space: $\mathbf{x} = (x_0, x_1, x_2)^T$
 - x_0, x_1 -> a point on the projection plane
 - synonyms: $(\text{tbf}\{\mathbf{x}\}, (\text{tbf}\{\mathbf{x}\}, x_2)^T, (x_0, x_1, x_2)^T$
 - light intensity:

$$I_\lambda(\mathbf{x}) = \int_0^L c_\lambda(\mathbf{x}, \xi) f'_c(\mathbf{x}, \xi) e^{-\int_0^\xi f'_c(\mathbf{x}, \mu) d\mu} d\xi, \quad (7)$$

- extinction function in ray space (after perform transformation from source space / object space):

$$f'_c(\mathbf{x}) = f_c(\varphi^{-1}(\phi^{-1}(\mathbf{x}))) = \sum_k w_k r'_k(\mathbf{x}), \quad (8)$$

- where φ is the mapping from object space to camera space, ϕ is the mapping from camera space to ray space
- substituting (8) into (7) we obtain a reconstruction kernel call it g_c :

$$I(\mathbf{x}) = \sum_k w_k \left(\int_0^L c(\mathbf{x}, \xi) r'_k(\mathbf{x}, \xi) \prod_j e^{-w_j \int_0^\xi r'_j(\mathbf{x}, \mu) d\mu} d\xi \right), \quad (9)$$

- Assumption:
 - Usually, the reconstruction kernels $r'_k(\mathbf{x})$ have local support. The splatting approach assumes that these local support areas do not overlap along a ray \mathbf{x} and the reconstruction kernels are ordered front to back
 - emission coefficient is constant in the support of each reconstruction kernel along a ray
 - approximate the exponential function with the first two terms of its Taylor expansion, thus $e^{-x} \approx 1 - x$
 - ignore self-occlusion
- yielding:

$$g_c(\mathbf{x}) = \sum_k w_k c_k(\mathbf{x}) q_k(\mathbf{x}) \prod_{j=0}^{k-1} (1 - w_j q_j(\mathbf{x})), \quad (10)$$

where $q_k(\mathbf{x})$ denotes an integrated reconstruction kernel, hence:

$$q_k(\mathbf{x}) = \int_{\mathbb{R}} r'_k(\mathbf{x}, x_2) dx_2. \quad (11)$$

- $q_k(\mathbf{x})$: footprint function of each point on the screen is an ellipse whose shape and orientation depend on the point's properties and the viewing parameters